

WolfWR Database Management System

CSC 540 Project Report 3

Matt Farver, Nick Garner, Jonathan Nguyen, Wyatt Plaga

Project Assumptions and Modifications

1. Merchandise can only have one Discount active at a time.
2. All *quantity* of a given Merchandise from a given Transaction must be returned together. Returns of one out of x of the same Merchandise are not allowed.
3. *transactionID*, *staffID*, and *memberID* are unique across all stores.
4. A single Warehouse processes all Supplier shipments for the entire chain, and is represented as a subclass of Store (*storeID*=2000).
5. *rewardAmount* and *amountOwed* are additional attributes attached to Member and Supplier entities, respectively, to automatically keep track of the money to be paid to each as shipments and transactions occur.
6. Members can decline to give *phone* or *email* when signing up at a store, hence why these fields are allowed to be NULL.
7. Every Staff member must have a *phone* and *email*, hence why these fields are NOT NULL.
8. *products* has been changed to *productID* and added to the primary key for Transaction, such that each real-world transaction will be split across multiple tuples with one unique product per tuple.
9. *storeID* has been added to the primary key for Merchandise, such that each Store's inventory of an item is stored on a separate tuple. This allows us to track shipments from Warehouse to Store and from Store to Store.
10. Not all products are perishable, thus *expiration* is allowed to be NULL in Merchandise.

Report 1 Corrections

Page #12 - Item 5: Information Processing APIs

Correction: Missing two API calls for managing promotions and sales information for products

Create/Update a Merchandise Sale

Input:

- DiscountID
- productID
- Start Date (start date of sale)
- End Date (end date of sale)
- Price Reduction

Response:

- Returns confirmation

GET/View Discounts

Input:

- DiscountID (optional - if null, will return all discounted merchandise)

Response:

- Returns confirmation
- DiscountID, ProductID, Price Reduction, Start Date, End Date

Page 11 - Items 9 & 10: Local Relation Schema + Documentation

Billing Operator

Correction: Staff(staffID, name, age, address, title, phone, email, employmentTime)

Member(memberID, firstName, lastName, address, phone, email, level, activeStatus, rewardAmount)

Supplier(supplierID, name, phone, email, location, amountOwed)

Correction: Transaction(transactionID, storeID, cashierID, date, total, productID)

Views(memberID, staffID)

Pays(supplierID, staffID)


has(transactionID, memberID) [able to be combined with the relation for the transaction entity set as it is many-to-one]

The E/R diagram for the warehouse operator view contained five entity sets and four relationships. The relationship schema for each of the entity sets was created in a straightforward manner, in which the key was listed first and underlined, followed by all of the attributes associated with that entity group. The “has” relation, being many-to-one, need not be converted into a relationship. If this were done, the memberID key would be included in the “transaction” relation. The other three relations, none having any attributes associated with the relation themselves, were converted simply by including all of the keys of the entity sets which

they connected. The “null-value” approach was used for treatment of the subclasses. Since the “Billing Operator” has no attributes not shared by the “Staff” entity set, the two entity sets were converted to a single “Staff” relationship schema. **Correction: Documentation now matches the schema above, as we have added the schema for Staff. As Billing Operator has no unique attributes, it is represented within the Staff schema itself.**

Page 7 - Item 7: Local E/R Diagrams

Correction: Changed “products” attribute to “productID”. Each transaction is split into separate tuples, one for each product. productID has been added to the primary key for Transaction.



Report 2 Corrections

Page #5 - Item 2: Design Decisions For Global Database Schema

Correction: schema Reward(rewardAmount, level) not discussed

Member(memberID, level, email, firstName, lastName, phone, address, rewardAmount, activeStatus)

Normalized forms:

Member(memberID, activeStatus, email, firstName, lastName, phone, address, rewardAmount)

~~Reward(rewardAmount, level)~~

- Functional Dependencies
 - memberID → activeStatus, level, email, firstName, lastName, phone, address, rewardAmount
 - In 3NF since the left side of the relation (memberID) contains a key.
 - ~~○ rewardAmount → level~~
 - ~~■ In 3NF since the left side of the relation (rewardAmount) now contains a key.~~
 - Correction: rewardAmount does *not* functionally determine level. A platinum member can still have a rewardAmount of \$0.00, such as when first signing up, or at the start of the year after reward checks have been sent out and rewardAmount is reset for the new year. Thus the Member relation is already in 3NF because there is only one functional dependency, the left side of which is a key (memberID). Therefore, there is no need for a Reward relation.

Page #4-8 - Item 2: Design Decisions For Global Database Schema

Staff(staffID, name, age, address, title, phone, email, employmentTime)

- Keys
 - Correction: reasons for the choosing of keys not well explained
 - staffID
 - Each staff member will have a unique staffID identifier. *As each staff has a unique staffID it will allow us to grab the other attributes of the staff member and allow us to display a particular staff member with the command SELECT * FROM Staff WHERE staffID = ?;*
- NULL
 - none
- NOT NULL
 - staffID, name, age, address, title, phone, email, employmentTime

- Each staff member will have name, address, title, phone, and email when hired.
 - A staff member should never have an employmentTime less than 1 so when a staff member is hired it will be set to a default of 1.

- Referential Integrity
 - none

Member(memberID, level, email, firstName, lastName, phone, address, rewardAmount, activeStatus)

- Keys

Correction: reasons for the choosing of keys not well explained

- memberID
 - Each member will have a unique memberID identifier. As each member has a unique memberID it will allow us to grab the other attributes of the member and allow us to display a particular member with the command `SELECT * FROM Member WHERE memberID = ?;`
- NULL
 - phone, email
 - Phone and email could be NULL if the customer declines to provide them. Means that the customer does not have a phone or email associated with their account.
- NOT NULL
 - memberID, activeStatus, level, firstName, lastName, address, rewardAmount
 - Every member will have a firstName, lastName, phone, address, and level when entered into the database.
 - activeStatus will be set to true if active and false if inactive
 - A member cannot have a reward amount less than 0 so this is set to a default value of 0.

- Referential Integrity

- memberID
 - memberID refers to the memberID of the SignUp Table.

Store(storeID, managerID, phone, address)

- Keys

Correction: reasons for the choosing of keys not well explained

- storeID
 - Each store will have a unique storeID identifier. As each store has a unique storeID it will allow us to grab the other attributes of the store and allow us to display a particular store with the command `SELECT * FROM Store WHERE storeID = ?;`

- - managerID
 - Each manager will have a unique managerID identifier.
- NULL
 - none
- NOT NULL
 - storeID, managerID, phone, address
 - Each store will have a storeID, managerID, phone, and address assigned.
- Referential Integrity
 - managerID
 - managerID refers to the managerID of the Staff Table.

Supplier(supplierID, name, phone, location, amountOwed, email)

- Keys

Correction: reasons for the choosing of keys not well explained

 - supplierID
 - Each supplier will have a unique supplierID identifier. As each supplier has a unique supplierID it will allow us to grab the other attributes of the supplier and allow us to display a particular supplier with the command SELECT * FROM Supplier WHERE supplierID = ?;
- NULL
 - none
- NOT NULL
 - supplierID, name, phone, email, location, amountOwed
 - Every supplier will have a supplierID, name, phone, email, and location when entered into the database.
 - A supplier cannot have an amount owed less than 0 so it is set to a default value of 0 when a supplier is entered into the database for the first time.
- Referential Integrity
 - none

Merchandise(productID, supplierID, storeID, name, quantity, buyPrice, marketPrice, expiration, productionDate)

- Keys

Correction: reasons for the choosing of keys not well explained

 - productID
 - Each product will have a unique productID identifier.
 - supplierID
 - Each supplier will have a unique supplierID identifier
 - storeID

- Each store will have a unique storeID identifier
 - As there may be multiple products with the same identifier, in order for a merchandise to be unique it must also contain a primary key of a supplierID and storeID. Thus a merchandise is unique as a specific product is from a specific supplier and at a specific store. This will allow us to grab a unique merchandise with the command `SELECT * FROM Merchandise WHERE productID = ? AND supplierID = ? AND storeID = ?;`
-
- NULL
 - expiration
 - A product could not have an expiration date.
- NOT NULL
 - productID, supplierID, storeID, name, quantity, buyPrice, marketPrice, productionDate
 - Each merchandise will have productID, supplierID, storeID, name, quantity, buyPrice, marketPrice, and productionDate when entered into the database.
- Referential Integrity
 - supplierID
 - supplierID refers to the supplierID of the Supplier Table.
 - storeID
 - storeID refers to the storeID of the Store Table.

Discount(productID, startDate, endDate, priceReduction)

- Keys

Correction: reasons for the choosing of keys not well explained

 - productID
 - Each discount will have a unique productID identifier
 - startDate
 - Each discount will have a unique startDate identifier
 - endDate
 - Each discount will have a unique endDate identifier.
 - As there may be multiple discounts on the same product, in order for a discount to be unique it must also contain a startDate and endDate primary key. The user may find a specific discount with the command `SELECT * FROM Discount WHERE discountID=?;` with each discount being unique being determined by the startDate and endDate.
- NULL
 - none
- NOT NULL
 - productID, startDate, endDate, priceReduction

- Each discount will have a productID, start date, end date, and how much the product is reduced by.
 - Referential Integrity
 - productID
 - productID refers to the productID of the Merchandise Table
- Transaction(transactionID, productID, total, date, cashierID, memberID, storeID, quantity)
- Keys

Correction: reasons for the choosing of keys not well explained

 - transactionID
 - Each transaction will have a unique transactionID identifier.
 - productID
 - Each product will have a unique productID identifier.
 - In order for a Transaction tuple to be unique, it must contain a unique combination of transactionID and productID. This allows us to split up one real-world transaction into multiple tuples with one unique product in each tuple. This helps facilitate writes of transactions with multiple products and simplifies the program logic behind product returns. The user may find a specific transaction with the command `SELECT * FROM Transaction WHERE transactionID=?;`
 -
 - NULL
 - none
 - NOT NULL
 - transactionID, productID, total, date, cashierID, memberID, storeID, quantity
 - Each transaction record will have transactionID, productID, total cost, quantity purchased, and date of purchase.
 - The quantity for a purchase cannot be less than 1 so it will be defaulted to 1 when entered into the system.
 - Along with each transaction record, it will contain the identifier of cashier, member, and store where the transaction took place.
 - Referential Integrity
 - cashierID, memberID, storeID, productID
 - cashierID refers to the staffID in the Staff Table.
 - memberID refers to the memberID in the Member Table.
 - storeID refers to the storeID in the Store Table.
 - productID refers to the productID in the Merchandise Table.

Sign-up(memberID, storeID, staffID, signUpDate)

- Keys

Correction: reasons for the choosing of keys not well explained

- memberID, storeID, staffID
 - Each member will have a unique memberID identifier.
 - Each store will have a unique storeID identifier.
 - Each staff member will have a unique staffID identifier.
- In order for a SignUp to be unique, the member must have a unique memberID from a unique storeID and by a unique staff member. This will allow for future query operations to find member growth reports by certain date or store. Such as with the command `SELECT COUNT(memberID) FROM SignUp WHERE storeID = ? AND signUpDate BETWEEN ? AND ?;`
-
- NULL
 - none
- NOT NULL
 - memberID, storeID, staffID, signUpDate
 - Each sign up will contain a memberID, storeID, and signUpDate
 - A signUpDate will always exist when a member signs up to the wholesale chain.
- Referential Integrity
 - storeID, staffID
 - storeID refers to the storeID in the Store Table
 - staffID refers to the staffID in the Staff Table

Transactions in Code

Example 1: MerchandiseSQL.returnInventory - MerchandiseSQL.java Lines 316 - 407

```
MerchandiseSQL.java
316  /**
317   * Method is called in Merchandise file under returnInventory and is used to execute queries
    when a member returns a product. The first query will be updating the
318   * Merchandise database and increasing the quantity of the product. The second query will be
    reducing the reward amount that the member has.
319   * The last query will be removing the transaction from the database.
320   * @param productID
321   * @param supplierID
322   * @param storeID
323   * @param memberID
324   * @param transactionID
325   * @throws SQLException
326   * @throws ParseException
327   */
328  public static void returnInventory(int productID, int supplierID, int storeID, int memberID,
    int transactionID, int transactionQuantity) throws SQLException, ParseException{
329      //Object that represents a precompiled SQL statement
330      PreparedStatement ps = null;
331      PreparedStatement ps2 = null;
332      PreparedStatement ps3 = null;
333      PreparedStatement memberSearch = null;
334
335      ResultSet member = null;
336
337      int id = 0;
338      int id2 = 0;
339      int id3 = 0;
340
341      try{
342          // Check if member is Platinum level
343          memberSearch = connection.prepareStatement("SELECT * FROM Member WHERE memberID = ?");
344          memberSearch.setInt(1, memberID);
345          member = memberSearch.executeQuery();
346          member.next();
347          boolean isPlatinum = member.getString("level").toLowerCase().equals("platinum");
348
349          // Turn off auto-commit
350          connection.setAutoCommit(false);
351
352          ps = connection.prepareStatement("Update Merchandise SET quantity = quantity + ? WHERE
    productID = ? AND supplierID = ? AND storeID = ?");
353          ps.setInt(1, transactionQuantity);
354          ps.setInt(2, productID);
355          ps.setInt(3, supplierID);
356          ps.setInt(4, storeID);
357
358          id = ps.executeUpdate();
359          System.out.println(id);
360      }
```

```

360
361     if (isPlatinum) {
362         ps2 = connection.prepareStatement("UPDATE Member SET rewardAmount = rewardAMOUNT -
        (SELECT(total*0.02) FROM Transaction WHERE transactionID = ? AND productID = ?)
        WHERE memberID = ?;");
363         ps2.setInt(1,transactionID);
364         ps2.setInt(2,productID);
365         ps2.setInt(3,memberID);
366
367         id2 = ps2.executeUpdate();
368         System.out.println(id2);
369     }
370
371     ps3 = connection.prepareStatement("DELETE FROM Transaction WHERE transactionID = ?;");
372     ps3.setInt(1,transactionID);
373
374     id3 = ps3.executeUpdate();
375     System.out.println(id3);
376
377     connection.commit();
378
379     if(id > 0){
380         System.out.println("Merchandise quantity updated");
381     } else{
382         System.out.println("Merchandise quantity not updated");
383     }
384
385     if (isPlatinum) {
386         if(id2 > 0){
387             System.out.println("Member reward amount updated");
388         } else{
389             System.out.println("Member reward amount not updated");
390         }
391     }
392
393     if(id3 > 0){
394         System.out.println("Transaction deleted");
395     } else{
396         System.out.println("Transaction not deleted");
397     }
398
399     connection.setAutoCommit(true);
400 }
401
402 catch (SQLException e) {
403     System.out.println("SQL Exception");
404     connection.rollback();
405     e.printStackTrace();
406     connection.setAutoCommit(true);
407 }

```

Explanation: In this method we are taking the inputs gathered by the menu and processing the inventory return request. We wrap the entire SQL workflow in a try/catch block so that if at any point we encounter a SQLException we can rollback the transactions. We start the method with a standalone transaction that simply queries the Member relation to check if the member in question is platinum level. We do this so we can know if we need to adjust the member's

rewardAmount as a result of the return. After this check, we turn auto-commit off and group the rest of the SQL statements as one transaction. Statement ps updates the Merchandise table to return the quantity of the product in question to the store it was purchased from. Statement ps2 is used if the member is platinum level to deduct 2% of the purchase total from their rewardAmount. Finally, statement ps3 deletes the Transaction tuple associated with this product. We then execute the COMMIT statement on line 377. If we have made it this far without any SQL Exceptions then the appropriate print statements are output and the catch block is not triggered to rollback the transaction.

Example 2: MerchandiseSQL.transferInventory - MerchandiseSQL.java Lines 408 - 484

```
MerchandiseSQL.java
408  /**
409   * This method is called in the Merchandise file under transferInventory. The purpose of this
410   * method is to take in two storeIDs and transfer the product of one
411   * store to another store. This method will execute two queries with the first query reducing
412   * the quantity of the original store. The second query will be adding the
413   * product to the second store, while handling duplicates.
414   * @param productID
415   * @param storeID
416   * @param name
417   * @param quantity
418   * @param buyPrice
419   * @param marketPrice
420   * @param productionDate
421   * @param expiration
422   * @param supplierID
423   * @param storeID2
424   * @param xferQuantity quantity to transfer
425   * @throws SQLException
426   * @throws ParseException
427   */
428 public static void transferInventory(int productID, int storeID, String name, int quantity,
429 double buyPrice, double marketPrice, Date productionDate, Date expiration, int supplierID, int
430 storeID2, int xferQuantity) throws SQLException, ParseException{
431     //Object that represents a precompiled SQL statement
432     PreparedStatement ps = null;
433     PreparedStatement ps2 = null;
434     int id = -1;
435     int id2 = -1;
436
437     try{
438         // Turn off auto-commit
439         connection.setAutoCommit(false);
440
441         ps = connection.prepareStatement("UPDATE Merchandise SET quantity = quantity - ? WHERE
442 productID = ? AND supplierID = ? AND storeID = ?;");
443 ps.setInt(1, xferQuantity);
444 ps.setInt(2, productID);
445 ps.setInt(3, supplierID);
446 ps.setInt(4, storeID);
447
448 id = ps.executeUpdate();
449 System.out.println(id);
```

```

445
446     ps2 = connection.prepareStatement("INSERT INTO Merchandise (productID, storeID, name,
quantity, buyPrice, marketPrice, productionDate, expiration, supplierID) VALUES (?, ?, ?, ?
, ?, ?, ?, ?) ON DUPLICATE KEY UPDATE quantity = quantity + ?;");
447     ps2.setInt(1, productID);
448     ps2.setInt(2, storeID2);
449     ps2.setString(3, name);
450     ps2.setInt(4, xferQuantity);
451     ps2.setDouble(5, buyPrice);
452     ps2.setDouble(6, marketPrice);
453     ps2.setDate(7, productionDate);
454     ps2.setDate(8, expiration);
455     ps2.setInt(9, supplierID);
456     ps2.setInt(10, xferQuantity);
457
458     id2 = ps2.executeUpdate();
459     System.out.println(id2);
460
461     connection.commit();
462
463     if(id > 0){
464         System.out.println("Merchandise removed from original store");
465     } else{
466         System.out.println("Merchandise not removed from original store");
467     }
468
469     if(id2 > 0){
470         System.out.println("Merchandise transfered successfully to store");
471     } else{
472         System.out.println("Merchandise not transfered successfully to store");
473     }
474
475     connection.setAutoCommit(true);
476 }
477 catch (SQLException e) {
478     System.out.println("SQL Exception");
479     connection.rollback();
480     e.printStackTrace();
481     connection.setAutoCommit(true);
482 }
483 }
484 }

```

Explanation: Similar to returnInventory, this method takes user inputs from the menu and executes the SQL statements to transfer inventory between stores. We surround the entire method with a try/catch block to catch any SQL Exceptions and rollback the transaction. After turning auto-commit off, our first statement updates the Merchandise relation to remove the quantity being transferred from the sending Store. Our second statement Inserts the transferred product and quantity as a new tuple with the receiving Store, unless the receiving Store already has a record of this product, in which case the transferred quantity is simply added to the quantity of the existing tuple. After these two statements are built and sent, COMMIT is called on line 461 for the transaction. Assuming we have gotten to this point without any exceptions, our print statements are output and we turn auto-commit back on.

High Level Design

In planning the high-level design of our system, we first made the decision to use a menu-based scheme to prompt for and receive user input. This seemed like the best alternative to a traditional UI and allowed us to concentrate on the backend logic and JDBC implementation. The decision to use Java and JDBC to interface with MariaDB was made in part due to the teaching staff's recommendation and our team's familiarity with those tools. In structuring our files, we tried to keep our code aligned with the database relations on which each file was interacting with. For example, Merchandise operations would be implemented in Merchandise.java. For the sake of readability and conciseness we chose to implement menu and IO logic in <RelationName> files and SQL statements in <RelationSQL> files. In cases where an operation touched more than one relation, the methods were implemented in the file of the relation most heavily affected by the operation. For example, returnInventory updates both the Member and Merchandise relations, but aligns more with the Merchandise relation in both use case and number of database interactions.

Development work began with Create-Read-Update-Destroy (CRUD) code for all relations. This was blackbox tested before moving on. The CRUD code was then extended to account for the other operations in the narrative, adding new methods for some operations and extending existing methods for others. Some common exception types, mainly SQLException, are caught and handled by most of our methods but error checking is not very extensive beyond that, in accordance with the project assignment. We also made an effort to output helpful print statements to illustrate what is happening in the database during an operation, for example outputting success statements as well as statements that describe what values were changed and by what amount.

Team Organization

Report 1

- Database Designer
 - Nick Garner (prime)
 - Jonathan Nguyen (backup)
- Application Programmer
 - Matt Farver
- Test Plan Engineer
 - Jonathan Nguyen (prime)
 - Wyatt Plaga (backup)

Report 2

- Software Engineer
 - Nick Garner (prime)
 - Matt Farver (backup)
- Database Designer
 - Jonathan Nguyen (prime)
 - Wyatt Plaga (backup)

Report 3 and Code

- Software Engineer
 - Nick Garner (prime)
 - Jonathan Nguyen (backup)
- Database Administrator
 - Jonathan Nguyen (prime)
 - Nick Garner (backup)
- Application Programmer
 - Jonathan Nguyen (prime)
 - Nick Garner (backup)
 - Matt Farver (backup)
 - Wyatt Plaga (backup)
- Test Plan Engineer
 - Nick Garner (prime)
 - Jonathan Nguyen (backup)